

### Advance DevOps Lab Experiment 03

**Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.**

<b><u>Roll No.</u></b>	<b><u>22</u></b>
<b><u>Name</u></b>	<b><u>Sarthak Harade</u></b>
<b><u>Class</u></b>	<b><u>D15B</u></b>
<b><u>Subject</u></b>	<b><u>Advance DevOps Lab</u></b>
<b><u>LO Mapped</u></b>	<b><u>LO1: To understand the fundamentals of Cloud Computing and be fully proficient with Cloud based DevOps solution deployment options to meet your business requirements.</u></b> <b><u>LO2: To deploy single and multiple container applications and manage application deployments with rollouts in Kubernetes</u></b>
<b><u>Grade:</u></b>	

**Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.**

**Theory:**

Container-based microservices architectures have profoundly changed the way development and operations teams test and deploy modern software. Containers help companies modernize by making it easier to scale and deploy applications, but containers have also introduced new challenges and more complexity by creating an entirely new infrastructure ecosystem.

Large and small software companies alike are now deploying thousands of container instances daily, and that's a complexity of scale they have to manage. So how do they do it?

Enter the age of Kubernetes.

Originally developed by Google, Kubernetes is an open-source container orchestration platform designed to automate the deployment, scaling, and management of containerized applications. In fact, Kubernetes has established itself as the defacto standard for container orchestration and is the flagship project of the Cloud Native Computing Foundation (CNCF), backed by key players like Google, AWS, Microsoft, IBM, Intel, Cisco, and Red Hat.

Kubernetes makes it easy to deploy and operate applications in a microservice architecture. It does so by creating an abstraction layer on top of a group of hosts so that development teams can deploy their applications and let Kubernetes manage the following activities:

- Controlling resource consumption by application or team
- Evenly spreading application load across a hosting infrastructure
- Automatically load balancing requests across the different instances of an application
- Monitoring resource consumption and resource limits to automatically stop applications from consuming too many resources and restarting the applications again
- Moving an application instance from one host to another if there is a shortage of resources in a host, or if the host dies
- Automatically leveraging additional resources made available when a new host is added to the cluster
- Easily performing canary deployments and rollbacks

## Steps:-

## Security Groups:-

## Master:-

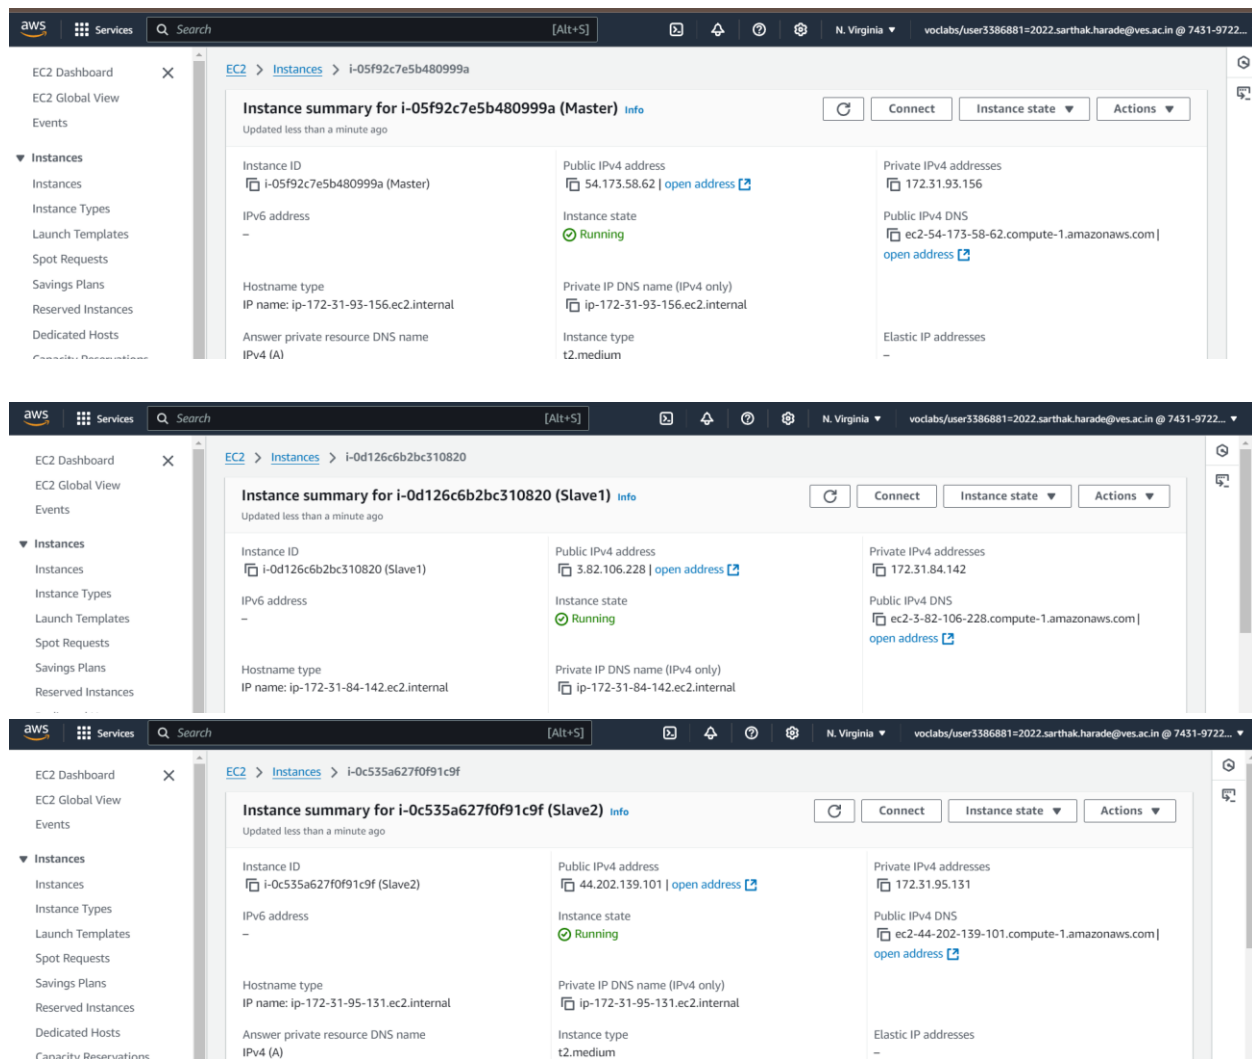
Inbound rules <a href="#">Info</a>							
Security group rule ID	Type <a href="#">Info</a>	Protocol <a href="#">Info</a>	Port range <a href="#">Info</a>	Source <a href="#">Info</a>	Description - optional <a href="#">Info</a>		
sgr-0c17c1a22a7c7b3e5	HTTP	TCP	80	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-0d3f86194443b29f1	All traffic	All	All	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-010d128b1484ff522	Custom TCP	TCP	6443	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-05bb413f0626b9c3b	Custom TCP	TCP	10251	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-04bd998c8f409420d	Custom TCP	TCP	10250	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-01438a40425cf867c	All TCP	TCP	0 - 65535	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-05dc20a8c2b541402	Custom TCP	TCP	10252	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-08d45afaf6c06c26	SSH	TCP	22	Custom	Q		Delete
					0.0.0.0/0 X		

## Slave:-

Security group rule ID							
Security group rule ID	Type <a href="#">Info</a>	Protocol <a href="#">Info</a>	Port range <a href="#">Info</a>	Source <a href="#">Info</a>	Description - optional <a href="#">Info</a>		
sgr-0dd83454961e5d3e9	All traffic	All	All	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-0402e9e84cd3dea45	SSH	TCP	22	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-05770af1e4c56697f	Custom TCP	TCP	10250	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-063fbc7516970bc90	All TCP	TCP	0 - 65535	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-07384bc31bec899e9	Custom TCP	TCP	30000 - 32767	Custom	Q		Delete
					0.0.0.0/0 X		
sgr-05188e46d7e21828d	HTTP	TCP	80	Custom	Q		Delete
					0.0.0.0/0 X		

## Creating 3 instances(1 Master, 2 Worker Nodes)(Use the same Key for connecting each instance through SSH to your command prompt)

Instances (4) <a href="#">Info</a>									
Last updated less than a minute ago									
Find Instance by attribute or tag (case-sensitive)									
	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4	
<input type="checkbox"/>	Master	i-05f92c7e5b480999a	Running	t2.medium	2/2 checks passed	View alarms +	us-east-1b	ec2-54-172...	
<input type="checkbox"/>	Master	i-03bb26146cdda2928	Terminated	t2.micro	-	View alarms +	us-east-1d	-	
<input type="checkbox"/>	Slave1	i-0d126c6b2bc310820	Running	t2.medium	2/2 checks passed	View alarms +	us-east-1b	ec2-3-82-1...	
<input type="checkbox"/>	Slave2	i-0c535a627f0f91c9f	Running	t2.medium	2/2 checks passed	View alarms +	us-east-1b	ec2-44-20...	



The image displays three screenshots of the AWS Management Console, showing the details of three EC2 instances in the N. Virginia region.

**Instance 1: Master**

- Instance ID: i-05f92c7e5b480999a (Master)
- Public IPv4 address: 54.173.58.62 | [open address](#)
- Private IPv4 addresses: 172.31.93.156
- Instance state: **Running**
- Public IPv4 DNS: ec2-54-173-58-62.compute-1.amazonaws.com | [open address](#)
- Private IP DNS name (IPv4 only): ip-172-31-93-156.ec2.internal
- Instance type: t2.medium

**Instance 2: Slave1**

- Instance ID: i-0d126c6b2bc310820 (Slave1)
- Public IPv4 address: 3.82.106.228 | [open address](#)
- Private IPv4 addresses: 172.31.84.142
- Instance state: **Running**
- Public IPv4 DNS: ec2-3-82-106-228.compute-1.amazonaws.com | [open address](#)
- Private IP DNS name (IPv4 only): ip-172-31-84-142.ec2.internal

**Instance 3: Slave2**

- Instance ID: i-0c535a627f0f91c9f (Slave2)
- Public IPv4 address: 44.202.139.101 | [open address](#)
- Private IPv4 addresses: 172.31.95.131
- Instance state: **Running**
- Public IPv4 DNS: ec2-44-202-139-101.compute-1.amazonaws.com | [open address](#)
- Private IP DNS name (IPv4 only): ip-172-31-95-131.ec2.internal
- Instance type: t2.medium

Connect to your EC2 instance through SSH by running the command having such format:-  
ssh -i "<Your\_saved\_key>.pem" ubuntu@<your-instance-public-ip>.<the region of created instance>.compute.amazonaws.com

Showed when you go to the SSH Client section when you select your instance and press Connect like here:-


EC2 Instance Connect

Session Manager


SSH client


EC2 serial console

Instance ID


 i-0af54010ae84808d2 (Node 2)


1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is Node1.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
 

 `chmod 400 "Node1.pem"`
4. Connect to your instance using its Public DNS:
 

 `ec2-34-224-169-38.compute-1.amazonaws.com`

Example:

 `ssh -i "Node1.pem" ubuntu@ec2-34-224-169-38.compute-1.amazonaws.com`

 **Note:** In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

```

C:\Users\mishr\Downloads>ssh -i "MasterExp3.pem" ubuntu@ec2-35-154-119-86.ap-south-1.compute.amazonaws.com
The authenticity of host 'ec2-35-154-119-86.ap-south-1.compute.amazonaws.com (35.154.119.86)' can't be established.
ED25519 key fingerprint is SHA256:kGQNJNTTMy98jwa1XzGUZCk8m274BhMAdkg4MiYxPY.
This key is not known by any other names
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-35-154-119-86.ap-south-1.compute.amazonaws.com' (ED25519) to the list of known hosts.
Welcome to Ubuntu 24.04 LTS (GNU/Linux 6.8.0-1012-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/pro

System information as of Tue Oct  1 13:57:27 UTC 2024

System load:  0.06               Processes:            105
Usage of /:   22.7% of 6.71GB    Users logged in:     0
Memory usage: 19%               IPv4 address for enx0: 172.31.44.130
Swap usage:   0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

```

**Step 4: Run on Master, Node 1, and Node 2 the below commands to install and setup Docker in Master, Node1, and Node2.**

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee
/etc/apt/trusted.gpg.d/docker.gpg > /dev/null
```

```
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu
$(lsb_release -cs) stable"
```

```

ubuntu@ip-172-31-44-130:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee
/etc/apt/trusted.gpg.d/docker.gpg > /dev/null
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu
$(lsb_release -cs) stable"
Warning: apt-key is deprecated. Manage keyring files in trusted.gpg.d instead (see apt-key(8)).
OK
-----BEGIN PGP PUBLIC KEY BLOCK-----

mQINBfit2ioBEADhWpZ8/wvZ6hUTiX0wQHXMALaFHCpH9hAtr4F1y2+OYdbtMuth
lqqwp028AqyY+PRfVMtSYMBjuQuu5byyKR018bqYhuS3jtgQmLjZ/bJvXqnmVXh
38UuLa+z077PxyxQhu5BbqntTPQMfiyqEiU+BKbq2WmANUKQf+1AmZY/Iru0Xbnq
L4C1+gJ8vfXmQ299npCaxEjaNRVYf0S8QcixNzHUyNb6emjLANyEVLZzeqo7XKL7
UrwV5inawTSzWNvtjEjj4nJL8NsLwscLPQUhTQ+7BbQXAwAmeHCUTQIvVWxqw0N
cmhh4HgeQscQHYgQJjDVfOY5MucvgLbIgCqfzAHW9jxmRL4qbMZj+b1XoePEtht
ku4bIQN1X5P07fNwzlgRL5Z4P0XDDZTLIQ/EL58j9kp4bnWRCJW0lya+f8ocodo
vZZ+Doi+fy4D5ZGrL4XecIQP/Lv5uFyf+kQtL/94VFYVJ0LeAv8W92KdgDkhTcTD
G7c0tIkVEKNUq48b3aQ64N0ZQW7fVjfoKwEzd0qPE72Pa45jrZzvUFxSpdiNk2tZ
XYukHjlxExEgBdC/J3cMMNRE1F4NCA3ApfV1Y7/hTeOnmDuDYwr9/obA8t016Yljj
q5rdkywPf4JF8mXUW5eCN1vAFHxeg9ZWemhBtQmGxXnw9M+z6hWwc6ahmwARAQAB
tCtEb2NrZXIgmUmVsZWZfZSAoQ0UgZGVhKSA8ZG9ja2VyQGVRvY2tldi5jb20+IQI3
BBMBGcgAhBQJYrefAAhsvBQsJCAcDBRUKCQgLBRYCAwEAAh4BAheAAAOJEI2BgDw0
v82IsskP/iQZo68fLDQmNvn8X5XTd6RRaUH33kXYXquT6NkHJciS7E2gTJmqvMqd
tI4mNYHCSEYxISqrcYV5YqX9P6+Ko+vozo4nseUQLPH/ATQ4qL0Zok+ljkg3Lgk
jonyUf9bwtWxPf05HC3GMHPhhcUSexCxQLQvnFWXD2sWLKivHp2fT8QbRGeZ+d3m

get:27 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates/multiverse amd64 Components [212 B]
get:28 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates/multiverse amd64 c-n-f Metadata [532 B]
get:29 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports/main amd64 Components [208 B]
get:30 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports/main amd64 c-n-f Metadata [112 B]
get:31 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports/universe amd64 Packages [10.6 kB]
get:32 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports/universe Translation-en [10.8 kB]
get:33 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports/universe amd64 Components [17.6 kB]
get:34 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports/universe amd64 c-n-f Metadata [1104 B]
get:35 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports/restricted amd64 Components [216 B]
get:36 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports/restricted amd64 c-n-f Metadata [116 B]
get:37 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports/multiverse amd64 Components [212 B]
get:38 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports/multiverse amd64 c-n-f Metadata [116 B]
get:39 http://security.ubuntu.com/ubuntu noble-security/main amd64 Packages [380 kB]
get:40 http://security.ubuntu.com/ubuntu noble-security/main Translation-en [83.1 kB]
get:41 http://security.ubuntu.com/ubuntu noble-security/main amd64 c-n-f Metadata [4576 B]
get:42 http://security.ubuntu.com/ubuntu noble-security/universe amd64 Packages [274 kB]
get:43 http://security.ubuntu.com/ubuntu noble-security/universe Translation-en [116 kB]
get:44 http://security.ubuntu.com/ubuntu noble-security/universe amd64 Components [8632 B]
get:45 http://security.ubuntu.com/ubuntu noble-security/universe amd64 c-n-f Metadata [10.4 kB]
get:46 http://security.ubuntu.com/ubuntu noble-security/restricted amd64 Packages [353 kB]
get:47 http://security.ubuntu.com/ubuntu noble-security/restricted Translation-en [68.1 kB]
get:48 http://security.ubuntu.com/ubuntu noble-security/restricted amd64 c-n-f Metadata [428 B]
get:49 http://security.ubuntu.com/ubuntu noble-security/multiverse amd64 Packages [10.9 kB]
get:50 http://security.ubuntu.com/ubuntu noble-security/multiverse Translation-en [2808 B]
get:51 http://security.ubuntu.com/ubuntu noble-security/multiverse amd64 Components [208 B]
get:52 http://security.ubuntu.com/ubuntu noble-security/multiverse amd64 c-n-f Metadata [344 B]
etched 29.1 MB in 6s (4834 kB/s)
Reading package lists... Done
/: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring
(trusted.gpg), see the DEPRECATION section in apt-key(8) for details.

```

sudo apt-get update

sudo apt-get install -y docker-ce-cli-containerd.io

```

ubuntu@ip-172-31-44-130:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee /usr/share/keyrings/docker
-archive-keyring.gpg > /dev/null
ubuntu@ip-172-31-44-130:~$ echo "deb [arch=amd64 signed-by=/usr/share/keyrings/docker-archive-keyring.gpg] https://downl
oad.docker.com/linux/ubuntu $(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
ubuntu@ip-172-31-44-130:~$ sudo apt update
sudo apt install docker-ce docker-ce-cli containerd.io
Hit:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Hit:5 http://security.ubuntu.com/ubuntu noble-security InRelease
Err:4 https://download.docker.com/linux/ubuntu noble InRelease
  The following signatures couldn't be verified because the public key is not available: NO_PUBKEY 7EA0A9C3F273FCD8
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
143 packages can be upgraded. Run 'apt list --upgradable' to see them.
W: An error occurred during the signature verification. The repository is not updated and the previous index files will
be used. GPG error: https://download.docker.com/linux/ubuntu noble InRelease: The following signatures couldn't be verif
ied because the public key is not available: NO_PUBKEY 7EA0A9C3F273FCD8
W: Failed to fetch https://download.docker.com/linux/ubuntu/dists/noble/InRelease The following signatures couldn't be
verified because the public key is not available: NO_PUBKEY 7EA0A9C3F273FCD8
W: Some index files failed to download. They have been ignored, or old ones used instead.
W: Target Packages (stable/binary-amd64/Packages) is configured multiple times in /etc/apt/sources.list.d/archive-uri-ht
Unpacking slirp4netns (1.2.1-1build2) ...
Setting up docker-buildx-plugin (0.17.1-1~ubuntu.24.04~noble) ...
Setting up containerd.io (1.7.22-1) ...
Created symlink /etc/systemd/system/multi-user.target.wants/containerd.service → /usr/lib/systemd/system/containerd.serv
ice.
Setting up docker-compose-plugin (2.29.7-1~ubuntu.24.04~noble) ...
Setting up libltdl7:amd64 (2.4.7-7build1) ...
Setting up docker-ce-cli (5:27.3.1-1~ubuntu.24.04~noble) ...
Setting up libslirp0:amd64 (4.7.0-1ubuntu3) ...
Setting up pigz (2.8-1) ...
Setting up docker-ce-rootless-extras (5:27.3.1-1~ubuntu.24.04~noble) ...
Setting up slirp4netns (1.2.1-1build2) ...
Setting up docker-ce (5:27.3.1-1~ubuntu.24.04~noble) ...
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.
Processing triggers for man-db (2.12.0-4build2) ...
Processing triggers for libc-bin (2.39-0ubuntu8.2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.

```

**sudo mkdir -p /etc/docker**

**cat <<EOF | sudo tee /etc/docker/daemon.json**

```

{
"exec-opts": ["native.cgroupdriver=systemd"]
}
EOF

```

And

**sudo systemctl enable docker**

**sudo systemctl daemon-reload**

## sudo systemctl restart docker

```
ubuntu@ip-172-31-44-130:~$ sudo mkdir -p /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
ubuntu@ip-172-31-44-130:~$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
Synchronizing state of docker.service with SysV service script with /usr/lib/systemd/systemd-sysv-install.
Executing: /usr/lib/systemd/systemd-sysv-install enable docker
ubuntu@ip-172-31-44-130:~$ curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
gpg: missing argument for option "-o"
-bash: /etc/apt/keyrings/kubernetes-apt-keyring.gpg: No such file or directory
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /
```

**Step 5: Run the below command to install Kubernetes. curl -fsSL**

```
https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee
/etc/apt/sources.list.d/kubernetes.list
```

**And**

**sudo apt-get update**

**sudo apt-get install -y kubelet kubeadm kubectl**

**sudo apt-mark hold kubelet kubeadm kubectl**

**If any errors faced here please refer to:-**

**<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>**

**And**

**<https://kubernetes.io/docs/tasks/tools/install-kubectl-linux/>**



```

ubuntu@ip-172-31-44-130:~$ curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
ubuntu@ip-172-31-44-130:~$ echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /
ubuntu@ip-172-31-44-130:~$ sudo apt-get update
sudo apt-get install -y kubelet kubeadm kubectl
sudo apt-mark hold kubelet kubeadm kubectl
Hit:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Hit:5 http://security.ubuntu.com/ubuntu noble-security InRelease
Get:6 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb InRelease [1186 B]
Get:7 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb Packages [4865 B]
Fetched 6051 B in 1s (7212 B/s)
Reading package lists... Done
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  conntrack cri-tools kubernetes-cni
The following NEW packages will be installed:
  conntrack cri-tools kubeadm kubectl kubelet kubernetes-cni
0 upgraded, 6 newly installed, 0 to remove and 143 not upgraded.
Need to get 87.4 MB of archives.
After this operation, 314 MB of additional disk space will be used.
Get:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 conntrack amd64 1:1.4.8-1ubuntu1 [37.9 kB]
Get:2 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb cri-tools 1.31.1-1.1 [15.7 MB]

Selecting previously unselected package kubectl.
Preparing to unpack .../3-kubectl_1.31.1-1.1_amd64.deb ...
Unpacking kubectl (1.31.1-1.1) ...
Selecting previously unselected package kubernetes-cni.
Preparing to unpack .../4-kubernetes-cni_1.5.1-1.1_amd64.deb ...
Unpacking kubernetes-cni (1.5.1-1.1) ...
Selecting previously unselected package kubelet.
Preparing to unpack .../5-kubelet_1.31.1-1.1_amd64.deb ...
Unpacking kubelet (1.31.1-1.1) ...
Setting up conntrack (1:1.4.8-1ubuntu1) ...
Setting up kubectl (1.31.1-1.1) ...
Setting up cri-tools (1.31.1-1.1) ...
Setting up kubernetes-cni (1.5.1-1.1) ...
Setting up kubeadm (1.31.1-1.1) ...
Setting up kubelet (1.31.1-1.1) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
kubelet set on hold.
kubeadm set on hold.

```

**sudo systemctl enable --now kubelet**

**sudo apt-get install -y containerd**

```
ubuntu@ip-172-31-44-130:~$ sudo systemctl enable --now kubelet
sudo apt-get install -y containerd
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-ce-rootless-extras libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following additional packages will be installed:
  runc
The following packages will be REMOVED:
  containerd.io docker-ce
The following NEW packages will be installed:
  containerd runc
0 upgraded, 2 newly installed, 2 to remove and 143 not upgraded.
Need to get 47.2 MB of archives.
After this operation, 53.1 MB disk space will be freed.
Get:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 runc amd64 1.1.12-0ubuntu3.1 [8599 kB]
Get:2 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 containerd amd64 1.7.12-0ubuntu4.1 [38.6 MB]
Fetched 47.2 MB in 1s (63.8 MB/s)
(Reading database ... 68064 files and directories currently installed.)
Removing docker-ce (5:27.3.1-1~ubuntu.24.04~noble) ...
Removing containerd.io (1.7.22-1) ...
Selecting previously unselected package runc.
(Reading database ... 68044 files and directories currently installed.)
Preparing to unpack .../runc_1.1.12-0ubuntu3.1_amd64.deb ...
Unpacking runc (1.1.12-0ubuntu3.1) ...
Selecting previously unselected package containerd.
Preparing to unpack .../containerd_1.7.12-0ubuntu4.1_amd64.deb ...
```

**sudo mkdir -p /etc/containerd**

**sudo containerd config default | sudo tee /etc/containerd/config.toml**

**And**

**sudo systemctl restart containerd**

**sudo systemctl enable containerd**

**sudo systemctl status containerd**

```
ubuntu@ip-172-31-44-130:~$ sudo mkdir -p /etc/containerd
sudo containerd config default | sudo tee /etc/containerd/config.toml
disabled_plugins = []
imports = []
oom_score = 0
plugin_dir = ""
required_plugins = []
root = "/var/lib/containerd"
state = "/run/containerd"
temp = ""
version = 2

[cgroup]
  path = ""

[debug]
  address = ""
  format = ""
  gid = 0
  level = ""
  uid = 0

[grpc]
  address = "/run/containerd/containerd.sock"
  gid = 0
  max_recv_message_size = 16777216
  max_send_message_size = 16777216
  tcp_address = ""
  tcp_tls_ca = ""
```

```

sudo systemctl enable containerd
sudo systemctl status containerd
● containerd.service - containerd container runtime
   Loaded: loaded (/usr/lib/systemd/system/containerd.service; enabled; preset: enabled)
   Active: active (running) since Tue 2024-10-01 14:10:55 UTC; 417ms ago
     Docs: https://containerd.io
   Main PID: 5175 (containerd)
      Tasks: 6
   Memory: 16.3M (peak: 16.9M)
      CPU: 80ms
   CGroup: /system.slice/containerd.service
           └─5175 /usr/bin/containerd

Oct 01 14:10:55 ip-172-31-44-130 containerd[5175]: time="2024-10-01T14:10:55.330219996Z" level=info msg=serving... addr=
Oct 01 14:10:55 ip-172-31-44-130 containerd[5175]: time="2024-10-01T14:10:55.330271786Z" level=info msg=serving... addr=
Oct 01 14:10:55 ip-172-31-44-130 containerd[5175]: time="2024-10-01T14:10:55.330364943Z" level=info msg="Start subscrib
Oct 01 14:10:55 ip-172-31-44-130 containerd[5175]: time="2024-10-01T14:10:55.330408709Z" level=info msg="Start recoveri
Oct 01 14:10:55 ip-172-31-44-130 containerd[5175]: time="2024-10-01T14:10:55.330465552Z" level=info msg="Start event mo
Oct 01 14:10:55 ip-172-31-44-130 containerd[5175]: time="2024-10-01T14:10:55.330479955Z" level=info msg="Start snapshot
Oct 01 14:10:55 ip-172-31-44-130 containerd[5175]: time="2024-10-01T14:10:55.330498567Z" level=info msg="Start cni netw
Oct 01 14:10:55 ip-172-31-44-130 containerd[5175]: time="2024-10-01T14:10:55.330528137Z" level=info msg="Start streamin
Oct 01 14:10:55 ip-172-31-44-130 systemd[1]: Started containerd.service - containerd container runtime.
Oct 01 14:10:55 ip-172-31-44-130 containerd[5175]: time="2024-10-01T14:10:55.333274244Z" level=info msg="containerd suc
lines 1-21/21 (END)client_loop: send disconnect: Connection reset

```

## sudo apt-get install -y socat

```

ubuntu@ip-172-31-44-130:~$ sudo apt-get install -y socat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-ce-rootless-extras libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  socat
0 upgraded, 1 newly installed, 0 to remove and 143 not upgraded.
Need to get 374 kB of archives.
After this operation, 1649 kB of additional disk space will be used.
Get:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 socat amd64 1.8.0.0-4build3 [374 kB]
Fetched 374 kB in 0s (15.7 MB/s)
Selecting previously unselected package socat.
(Reading database ... 68108 files and directories currently installed.)
Preparing to unpack .../socat_1.8.0.0-4build3_amd64.deb ...
Unpacking socat (1.8.0.0-4build3) ...
Setting up socat (1.8.0.0-4build3) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

```

**Step 6: Initialize the Kubeccluster .Now Perform this Command only for Master. sudo kubeadm init --pod-network-cidr=10.244.0.0/16**

If any errors in this command , run:-

1.) Enable IP Forwarding by running the following command:

```
sudo sysctl -w net.ipv4.ip_forward=1
```

2.) Make the Change Persistent (to ensure it remains active after a reboot):

Open the system control configuration file for editing:

```
sudo nano /etc/sysctl.conf
```

3.) Add or modify the following line:

```
net.ipv4.ip_forward=1
```

4.) Save the file and exit the editor.(Click ctrl+X,then Yes and then Enter)

5.) Apply the Changes:

```
sudo sysctl -p
```

6.) Re-run the kubeadm init command:

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16
```

```
ubuntu@ip-172-31-44-130:~$ sudo sysctl -w net.ipv4.ip_forward=1

ubuntu@ip-172-31-44-130:~$ sudo sysctl -p
net.ipv4.ip_forward = 1
ubuntu@ip-172-31-44-130:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W1001 14:48:41.493013 1174 checks.go:846] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container runtime is inconsistent with that used by kubeadm.It is recommended to use "registry.k8s.io/pause:3.10" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-44-130.kubernetes.kubernetes.default.kubernetes.default.svc.kubernetes.default.svc.cluster.local] and IPs [10.96.0.1 172.31.44.130]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-44-130 localhost] and IPs [172.31.44.130 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-44-130 localhost] and IPs [172.31.44.130 127.0.0.1 ::1]
```

Run this command on master and also copy and save the Join command from below.

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.44.130:6443 --token bo0l2n.d8yziyz741mv80us \
  --discovery-token-ca-cert-hash sha256:c44dd007e132a163eed321a3d3dde1656aaedff3c30a5c42a563ffab6b7e02cd
ubuntu@ip-172-31-44-130:~$ mkdir -p $HOME/.kube
ubuntu@ip-172-31-44-130:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
ubuntu@ip-172-31-44-130:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
ubuntu@ip-172-31-44-130:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
ip-172-31-44-130    NotReady control-plane  44s   v1.31.1
```

Join command:-

```
kubeadm join 172.31.44.130:6443 --token bo0l2n.d8yziyz741mv80us \
  --discovery-token-ca-cert-hash sha256:c44dd007e132a163eed321a3d3dde1656aaedff3c30a5c42a563ffab6b7e02cd
```

**Step 8: Now Run the following command on Node 1 and Node 2 to Join to master. sudo kubeadm join <EC2 Instance Ip> --token <randomly\_alloted\_token>\ --discovery-token-ca-cert-hash**

**sha256:d6fc5fb7e984c83e2807780047fec6c4f2acfe9da9184ecc028d77157608fbb6**

```
ubuntu@ip-172-31-45-50:~$ sudo sysctl -p
net.ipv4.ip_forward = 1
ubuntu@ip-172-31-45-50:~$ sudo kubeadm join 172.31.44.130:6443 --token bo0l2n.d8yziyz741mv80us \
--discovery-token-ca-cert-hash sha256:c44dd007e132a163eed321a3d3dde1656aaedff3c30a5c42a563ffab6b7e02cd
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.001776345s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

```
ubuntu@ip-172-31-35-36:~$ sudo sysctl -w net.ipv4.ip_forward=1
net.ipv4.ip_forward = 1
ubuntu@ip-172-31-35-36:~$ sudo nano /etc/sysctl.conf
ubuntu@ip-172-31-35-36:~$ sudo sysctl -p
ubuntu@ip-172-31-35-36:~$ sudo nano /etc/sysctl.conf
ubuntu@ip-172-31-35-36:~$ sudo sysctl -p
net.ipv4.ip_forward = 1
ubuntu@ip-172-31-35-36:~$ sudo kubeadm join 172.31.44.130:6443 --token bo0l2n.d8yziyz741mv80us \
--discovery-token-ca-cert-hash sha256:c44dd007e132a163eed321a3d3dde1656aaedff3c30a5c42a563ffab6b7e02cd
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.001764003s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

**Step 9: Now Run the command kubectl get nodes to see the nodes after executing Join command on nodes.**

**And**

**Step 10: Since Status is NotReady we have to add a network plugin. And also we have to give the name to the nodes. kubectl apply -f**

**<https://docs.projectcalico.org/manifests/calico.yaml>**

```

ubuntu@ip-172-31-44-130:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE      VERSION
ip-172-31-35-36     NotReady <none>   3m52s    v1.31.1
ip-172-31-44-130    NotReady control-plane 8m36s    v1.31.1
ip-172-31-45-50     NotReady <none>   47s      v1.31.1
ubuntu@ip-172-31-44-130:~$ kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
poddisruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created

```

Now Run command kubectl get nodes -o wide

Or

kubectl get nodes so

we can see Status is ready.

Renaming:-

Rename to Node 1:kubectl label node ip-172-31-45-50 kubernetes.io/role=Node1

Rename to Node 2:kubectl label node ip-172-31-35-36 kubernetes.io/role=Node2

```

ubuntu@ip-172-31-44-130:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE      VERSION
ip-172-31-35-36     Ready     <none>   6m31s    v1.31.1
ip-172-31-44-130    Ready     control-plane 11m      v1.31.1
ip-172-31-45-50     Ready     <none>   3m26s    v1.31.1
ubuntu@ip-172-31-44-130:~$ kubectl label node ip-172-31-45-50 kubernetes.io/role=worker
node/ip-172-31-45-50 labeled
ubuntu@ip-172-31-44-130:~$ kubectl label node ip-172-31-45-50 kubernetes.io/role=Node1
error: 'kubernetes.io/role' already has a value (worker), and --overwrite is false
ubuntu@ip-172-31-44-130:~$ 2:kubectl label node ip-172-31-35-36 kubernetes.io/role=Node1
2:kubectl: command not found
ubuntu@ip-172-31-44-130:~$ kubectl label node ip-172-31-35-36 kubernetes.io/role=Node2
node/ip-172-31-35-36 labeled
ubuntu@ip-172-31-44-130:~$ kubectl get nodes -o wide
NAME                STATUS    ROLES    AGE      VERSION    INTERNAL-IP    EXTERNAL-IP    OS-IMAGE             KERNEL-VE
SION    CONTAINER-RUNTIME
ip-172-31-35-36     Ready     Node2    9m46s    v1.31.1    172.31.35.36   <none>         Ubuntu 24.04 LTS     6.8.0-101
2-aws   containerd://1.7.12
ip-172-31-44-130    Ready     control-plane 14m      v1.31.1    172.31.44.130 <none>         Ubuntu 24.04 LTS     6.8.0-101
2-aws   containerd://1.7.12
ip-172-31-45-50     Ready     worker   6m41s    v1.31.1    172.31.45.50   <none>         Ubuntu 24.04 LTS     6.8.0-101
2-aws   containerd://1.7.12

```

run kubectl get nodes

```

ubuntu@ip-172-31-44-130:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE      VERSION
ip-172-31-35-36     Ready     Node2    9m58s    v1.31.1
ip-172-31-44-130    Ready     control-plane 14m      v1.31.1
ip-172-31-45-50     Ready     worker   6m53s    v1.31.1
ubuntu@ip-172-31-44-130:~$ client_loop: send disconnect: Connection reset

```

**Conclusion:**

In this experiment, we learned how to install Kubernetes create a Kubernetes Cluster in AWS EC2 instances and get them up and running.